

# Sentiment Analysis

Human Language from a Computational Perspective  
June 27, 2018

# Outline

- Text classification
- Machine learning algorithms

# Text classification

Many problems involve **classification**:

- Topic [NEWS, SPORTS, ...]
- Author/gender [MALE, FEMALE]
- Spam filtering [SPAM, NOT-SPAM]

# Sentiment classification

Given an input sentence, return its  
**sentiment label:**

**+1 (POSITIVE)**   0 (NEUTRAL)   **-1 (NEGATIVE)**

(or a number on a finer scale, e.g. 1-5)

# News and stocks



# Twitter

[V\\_FSD](#): RT @FreeMemesKids: Love my new **fidget spinner** <https://t.co/lcA6Ui6fsV>

[MeqzEdits](#): @23Duckk Oh no the **fidget spinner**?

[fffaraa](#): @ibrahimyussop Omg.. not the **fidget spinner**

[boybrunch](#): RT @a1andar: The Born This Way album cover except Gaga is a **fidget spinner**

[eve\\_bertie](#): I bought a **fidget spinner** and I think it's the best thing I've ever done

[Hamza0207](#): I broke this little girls **fidget spinner** and I legit feel so bad ??

# Algorithm requirements

Input: sequence of tokens

{OH, NO, THE, FIDGET, SPINNER, ?}

Output: label (number)

-1

# Bag of words

Simple approach: ignore the order, and just look for indicative tokens:

NOT    GREAT    DAMN    LOVE    HATE    ...

Assign a weight to each token:

0            +1            -1            +2            -2            ...



# Bag of words: classification

Calculate score for sentence:

THE ONLY REASON I LOVE MONDAYS ... ZUMBA !!!!

$$0 + 0 + 0 + 0 + 0 + (+2) + (-1) + 0 + (+1) + (+1)$$

$$= +3$$

Since  $3 > 0$ , predict:

**+1 (POSITIVE)**

# Bag of words algorithm

**classify**(L, W):

$s \leftarrow 0$

$i \leftarrow 1$

while  $i \leq \text{len}(L)$ :

$s \leftarrow s + W[L[i]]$

$i \leftarrow i + 1$

return  $\text{sign}(s)$

▷ L: input sentence,

W: table of weights for words

+1 if  $s > 0$

0 if  $s = 0$

-1 if  $s < 0$

# Learning a model

How to determine the word weights?

In language models and POS tagging,  
we **learned** the statistics as **Counts**.

Here we can learn **Weights**.

# Perceptron

The perceptron is a **learning** algorithm.

Input: list of samples (sentence + label)

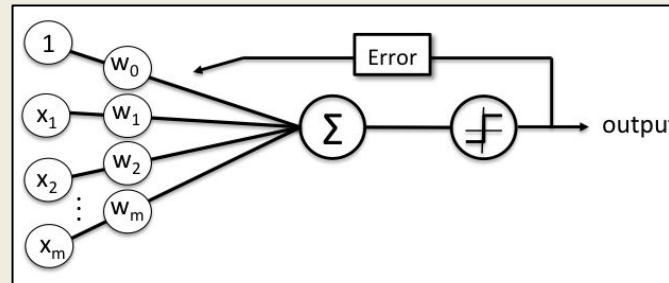
Labels given as numbers: **+1**    0    **-1**

Output: table of weights for each word

(which can then be used to classify new sentences.)

# Perceptron

The algorithm goes over all samples repeatedly, until there are no errors. Whenever there is an error, it updates the weights of all tokens in the sample.



# Perceptron algorithm

**train**(X, Y):

$W \leftarrow [0 \text{ for all words}]$

while W is changing:

$i \leftarrow 1$

while  $i \leq \text{len}(X)$ :

if **classify**(X[i], W)  $\neq$  Y[i]:

**update**(X[i], Y[i], W)

$i \leftarrow i + 1$

return W

▷ X: list of input sentence (samples)

Y: list of labels (one per sample)

▷ training iterations

▷ go through all samples

▷ if the model is wrong,  
update its weights

▷ return final learned weights

# Perceptron algorithm

**update(L, y, W):**

▷ L: input sentence, y: label,

$i \leftarrow 1$

W: table of weights for words

**while**  $i \leq \text{len}(L)$ :

▷ go through sentence tokens

$W[L[i]] \leftarrow W[L[i]] + y$

▷ update weight by

$i \leftarrow i + 1$

adding y to it

# Example

$x = \begin{bmatrix} \text{THIS IS GREAT,} \\ \text{THIS IS AWFUL} \end{bmatrix}$

$y = \begin{bmatrix} 1, \\ -1 \end{bmatrix}$



# Example

First iteration, first sample

$$X[1] = \begin{array}{l} \text{THIS IS GREAT} \\ 0 + 0 + 0 = 0 \end{array}$$

Result: 0 (NEUTRAL)

$$Y[1] = 1 \neq 0 \quad \text{X}$$

W =

THIS	0
IS	0
GREAT	0
AWFUL	0

# Example

First iteration, first sample

$$X[1] = \begin{matrix} \text{THIS IS GREAT} \\ 0 + 0 + 0 = 0 \end{matrix}$$


Result: 0 (NEUTRAL)

$$Y[1] = 1 \neq 0 \quad \text{X}$$

Updating weights

W =

THIS	1
IS	1
GREAT	1
AWFUL	0



# Example

First iteration, second sample

$$X[2] = \begin{matrix} \text{THIS IS AWFUL} \\ 1 + 1 + 0 = 0 \end{matrix}$$

Result: 1 (POSITIVE)

$$Y[2] = -1 \neq 1 \quad \text{X}$$

W =

THIS	1
IS	1
GREAT	1
AWFUL	0

# Example

First iteration, second sample

$$X[2] = \begin{matrix} \text{THIS IS AWFUL} \\ 1 + 1 + 0 = 0 \end{matrix}$$


Result: 1 (POSITIVE)

$$Y[2] = -1 \neq 1 \quad \text{X}$$

Updating weights

W =

THIS	0
IS	0
GREAT	1
AWFUL	-1



# Example

Second iteration, first sample

$$X[1] = \begin{matrix} \text{THIS IS GREAT} \\ 0 + 0 + 1 = 1 \end{matrix}$$

Result: 1 (POSITIVE)

$$Y[1] = 1$$



W =

THIS	0
IS	0
GREAT	1
AWFUL	-1

# Example

Second iteration, second sample

$$X[2] = \begin{array}{l} \text{THIS IS AWFUL} \\ 0 + 0 + \textcolor{red}{-1} = \textcolor{red}{-1} \end{array}$$

Result:  $\textcolor{red}{-1}$  (NEGATIVE)

$$Y[2] = \textcolor{red}{-1}$$



W =

THIS	0
IS	0
GREAT	1
AWFUL	-1

# Example

Finished iteration without updating weights, so return  $W$ .

$W$  can now be used to classify new sentences.

$W =$

THIS	0
IS	0
GREAT	1
AWFUL	-1

# Problem: order matters

THAT IS NOT TRUE , IT IS GREAT

has the same tokens as

THAT IS TRUE , IT IS NOT GREAT

Will be classified the same



# Generalizing

Instead of single words, use features.

Examples for features:

- first token is GREAT?
- contains NOT followed by BAD?
- contains TASTE with part of speech NN?

# Generalizing

In Bag of Words, all features are:

- contains  $\langle w \rangle$ ?

where  $\langle w \rangle$  is some word.

This is also called **unigram features**.

# Generalizing

We can also use **bigram features**:

- contains  $\langle w_1 \rangle$  followed by  $\langle w_2 \rangle$ ?

where  $\langle w_1 \rangle$  and  $\langle w_2 \rangle$  are words. Or:

- contains  $\langle w \rangle$  with part-of-speech  $\langle p \rangle$ ?

etc.

# Linear classifier

```
classify(L, F, W):  
    s ← 0  
    i ← 1  
    while i ≤ len(F):  
        f ← F[i]  
        if f(L):  
            s ← s + W[f]  
        i ← i + 1  
    return sign(s)
```

‣ L: input sentence, F: features to check,  
W: table of weights for features

‣ go through possible features

‣ the feature f is true for L

‣ add the feature's weight to the score

‣ 1 if  $s > 0$ , 0 if  $s = 0$ , -1 if  $s < 0$

# Linear classifier

The same as the Bag of Words classifier, but using general features and not just words.

Might be even millions of features!

# Example

[MegzEdits](#): @23Duckk Oh no the **fidget spinner**?

The unigram features  $O_H$  and  $NO$  may not be clearly negative (weights will stay  $\sim 0$ ), but the bigram feature  $O_H NO$  is negative.

# Multi-class classifier

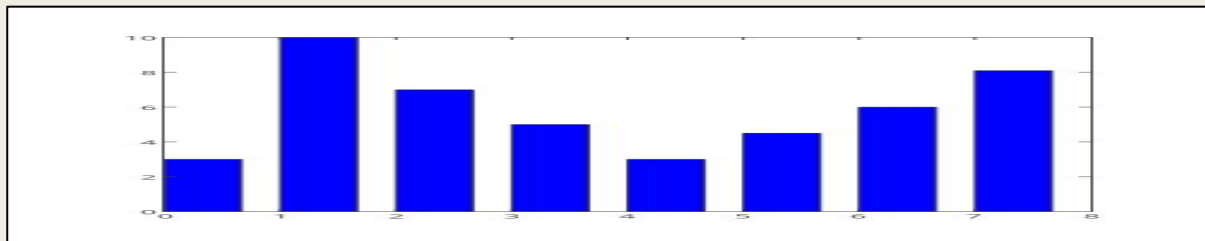
So far we discussed **binary** classifiers, but there can be more than two labels.

Example:

2 (VERY POSITIVE), 1 (POSITIVE), 0, -1 (NEGATIVE), -2 (VERY NEGATIVE)

# Multi-class classifier

Instead of a single score, the classifier will have a score for each possible label, and then choose the label with the top score.



Learning: different weights for each label.



# Linear multi-class classifier

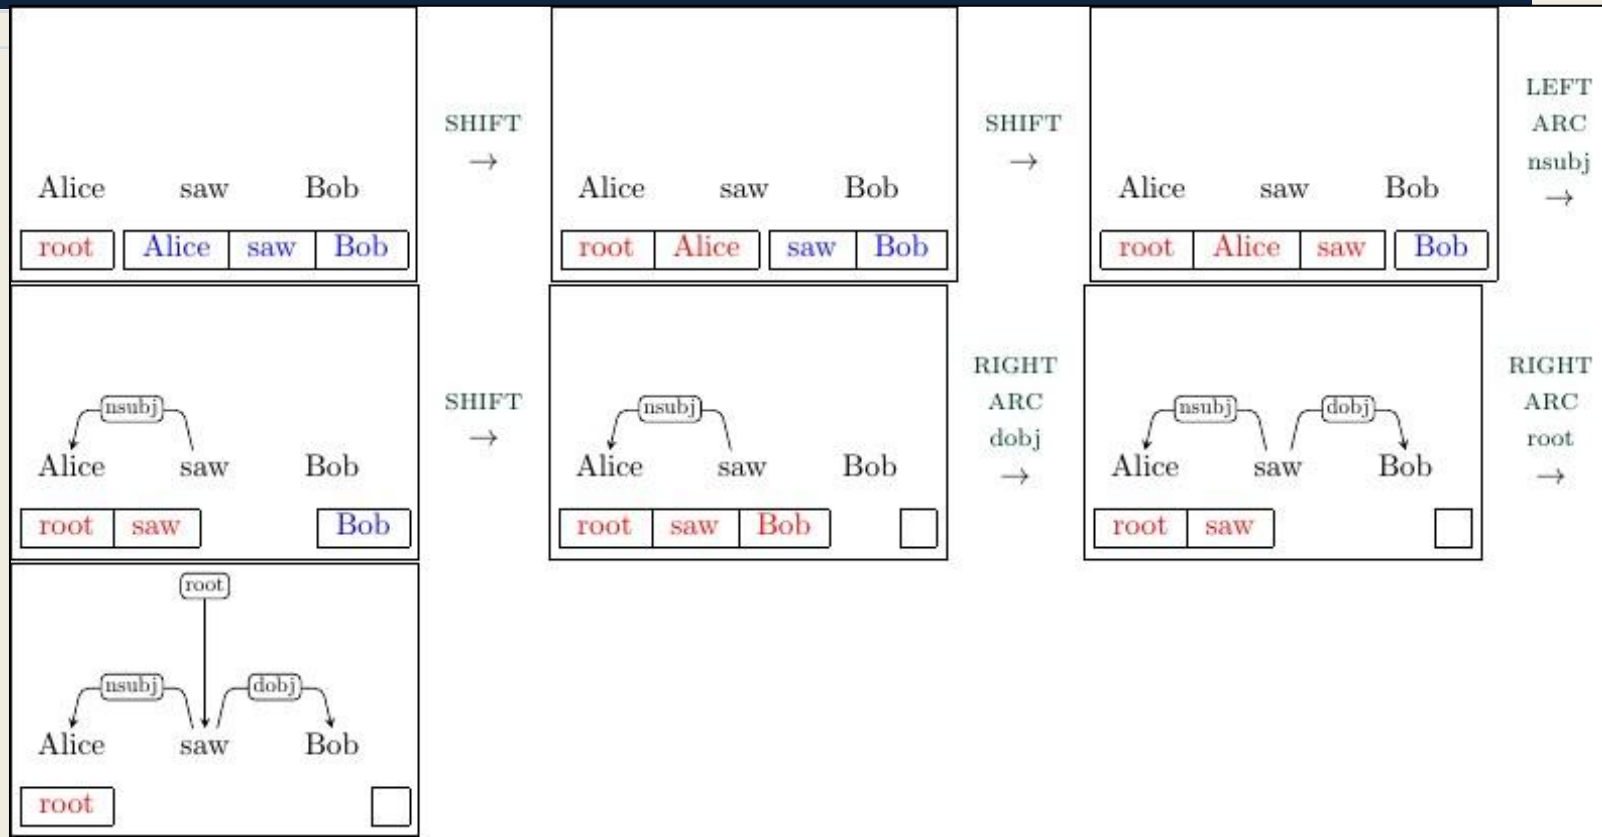
```

classify(L, F, W):
    S ← [0 for all labels]
    i ← 1
    while i ≤ len(F):
        f ← F[i]
        if f(L):
            S ← S + W[f]
        i ← i + 1
    return argmax(S)

```

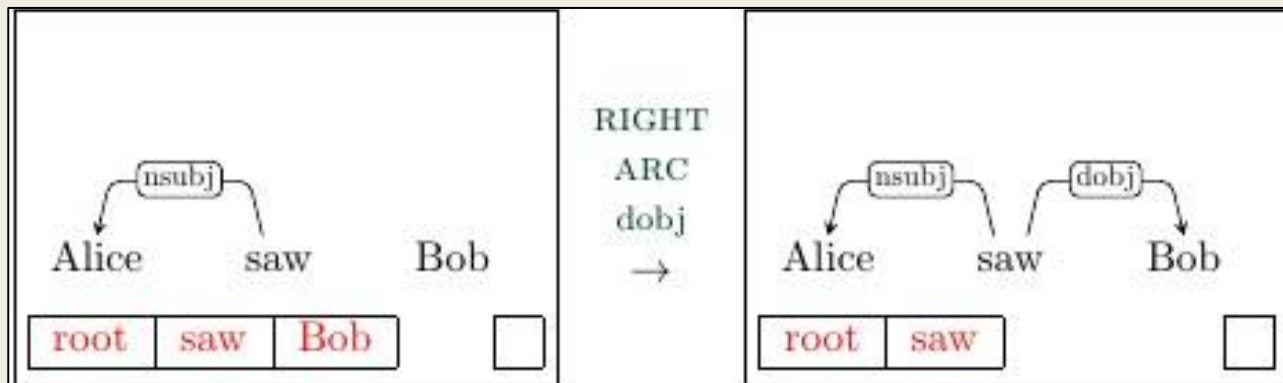
L: input sentence, F: features to check  
 W: table of weights for features for each label  
 go through possible features  
 the feature f is true for L  
 add the feature's weight to the score for each label

# Back to parsing



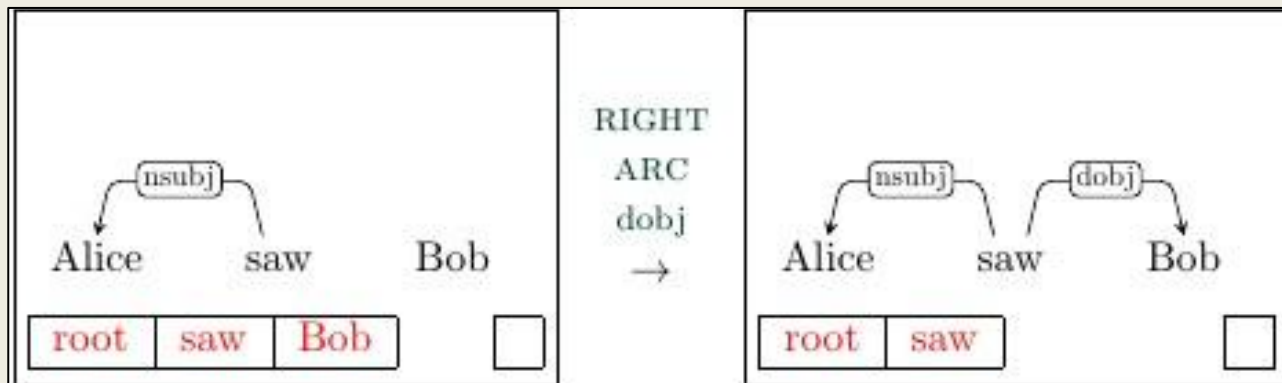
# Perceptron for parsing

If we see the state on the left here, we need to know to apply **RIGHT-ARC**<sub>dobj</sub>



# Features for parsing

- rightmost token on the stack is  $\langle w \rangle$ ?
- buffer leftmost part-of-speech is  $\langle P \rangle$ ?
- second stack token is a parent?
- etc.



# Perceptron for parsing

Perceptron + word/part-of-speech

unigram, bigram features: pretty good parser.

Parser	UAS (%)	LAS (%)
MaltParser	90.93	88.95
Parsey McParseface	94.41	92.55

# Twitter sentiment analysis

Training any model requires labeled data:  
learning from examples.



# Problem: labeling

Manually annotating tweets  
takes time and money.

# Avoiding manual work

Automatically building a training set using the **emoticons** in the tweets:

## POSITIVE

## NEGATIVE

:)	:D	: )	:(	=(	: (
:-)	=)	;)	:-( :/	:/	=/

Alec Go, Richa Bhayani, and Lei Huang. 2009. [Twitter sentiment classification using distant supervision](#).

Technical report, Stanford.



# Problem: sarcasm

People sometimes say the  
opposite of what they mean.

# Sarcasm recognition

Sarcasm: “saying the opposite of what you mean in a way intended to make someone else feel stupid or show you are angry”.

listening to Andrew Ridgley by Black Box Recorder on  
@Grooveshark: <http://tinysong.com/cO6i> #goodmusic

I guess you should expect a WONDERFUL video  
tomorrow. #sarcasm

Accuracy: 94.7%

1. Oren Tsur, Dmitry Davidov, Ari Rappoport. [ICWSM - A Great Catchy Name: Semi-Supervised Recognition of Sarcastic Sentences in Product Reviews. Fourth International AAAI Conference on Weblogs and Social Media \(ICWSM\) 2010.](#)
2. Oren Tsur, Dmitry Davidov, Ari Rappoport. [Semi-Supervised Recognition of Sarcastic Sentences in Twitter and Amazon. Computational Natural Language Learning \(CoNLL\) 2010.](#)

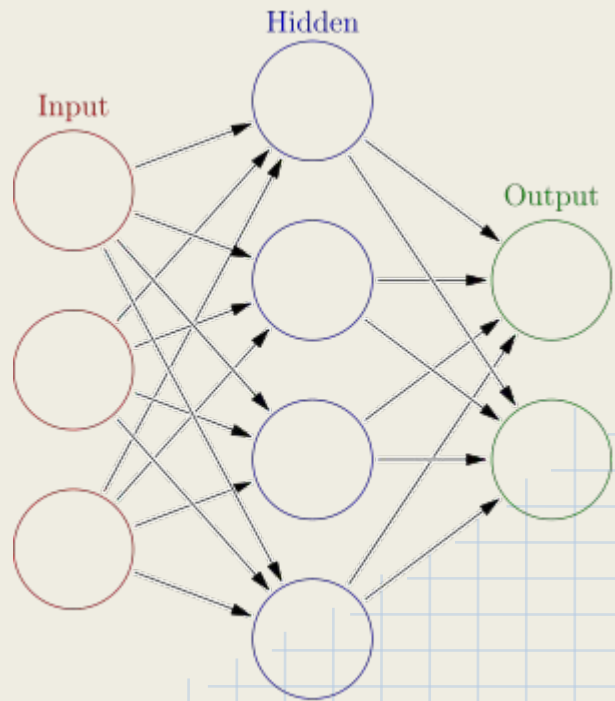
# Multi-class labels

Hashtag- and  
smiley-based labels:  
51 or 16 labels  
instead of just two  
(positive/negative)

Hashtags	Smileys
#sad	;)
#crazy	:(
#bored	X(
#fun	:S

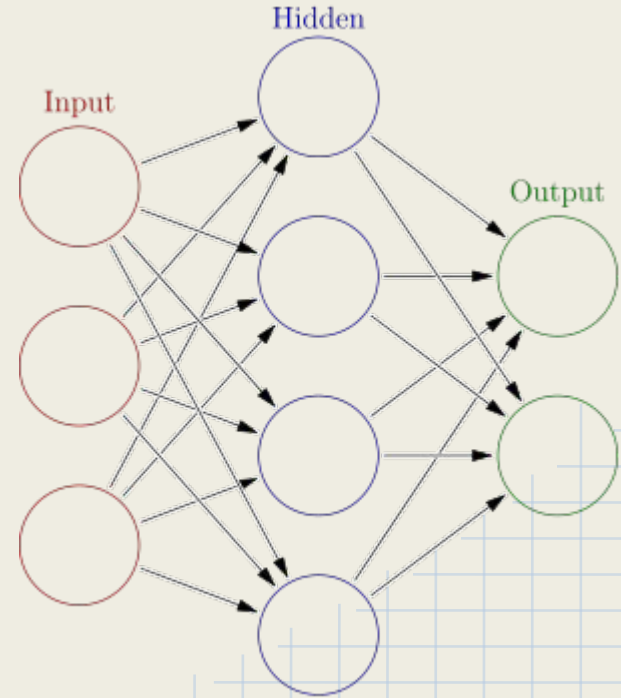
# Artificial neural networks

Instead of using the score for classification, give it to another linear classifier (“multi-layer perceptron”)



# Artificial neural networks

Only the output of the last layer is actually used for classification.

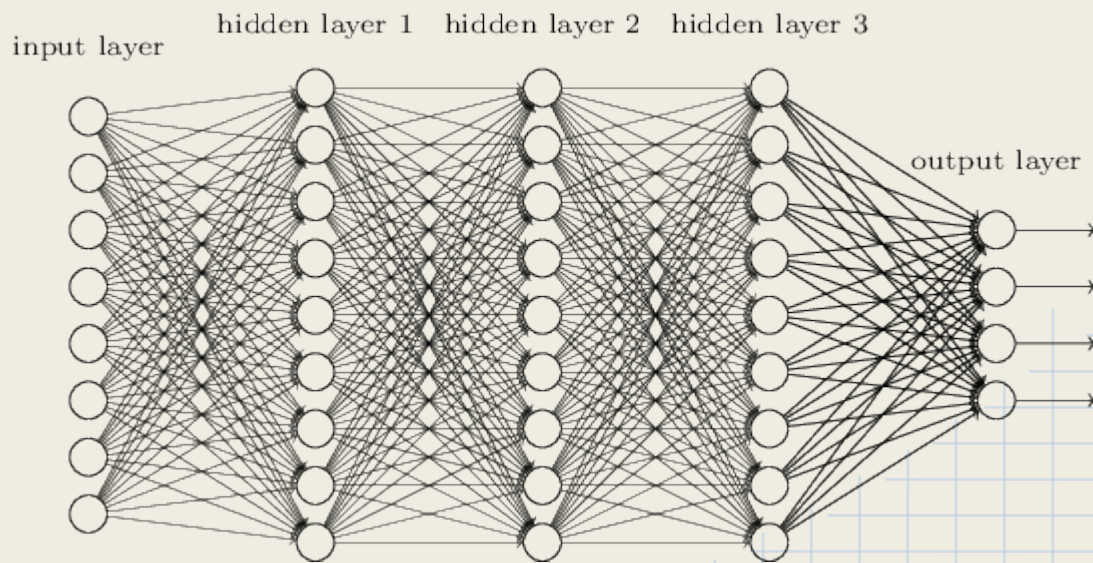


# Deep neural networks

Many layers

sometimes work  
better.

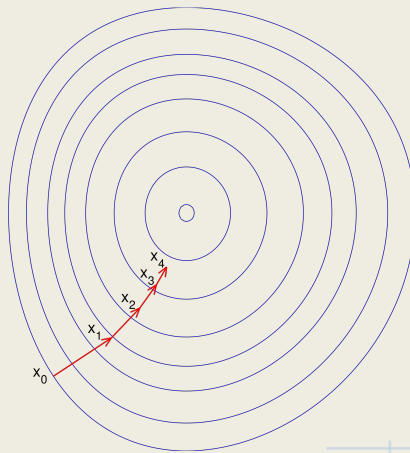
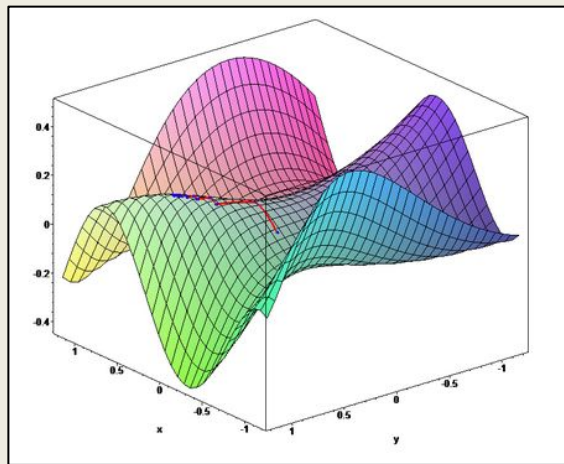
(“deep learning”)



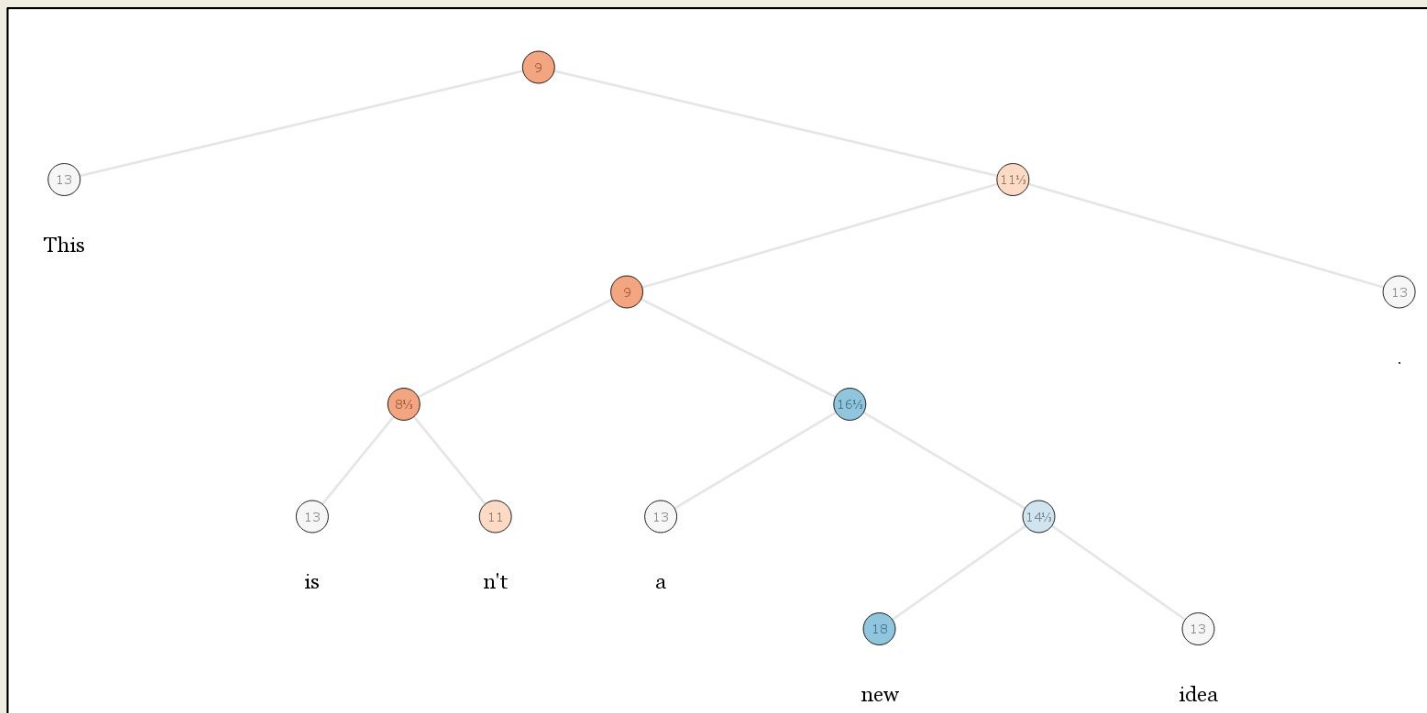
# Training neural networks

Perceptron only works for one layer.

Instead, use gradient descent algorithms.



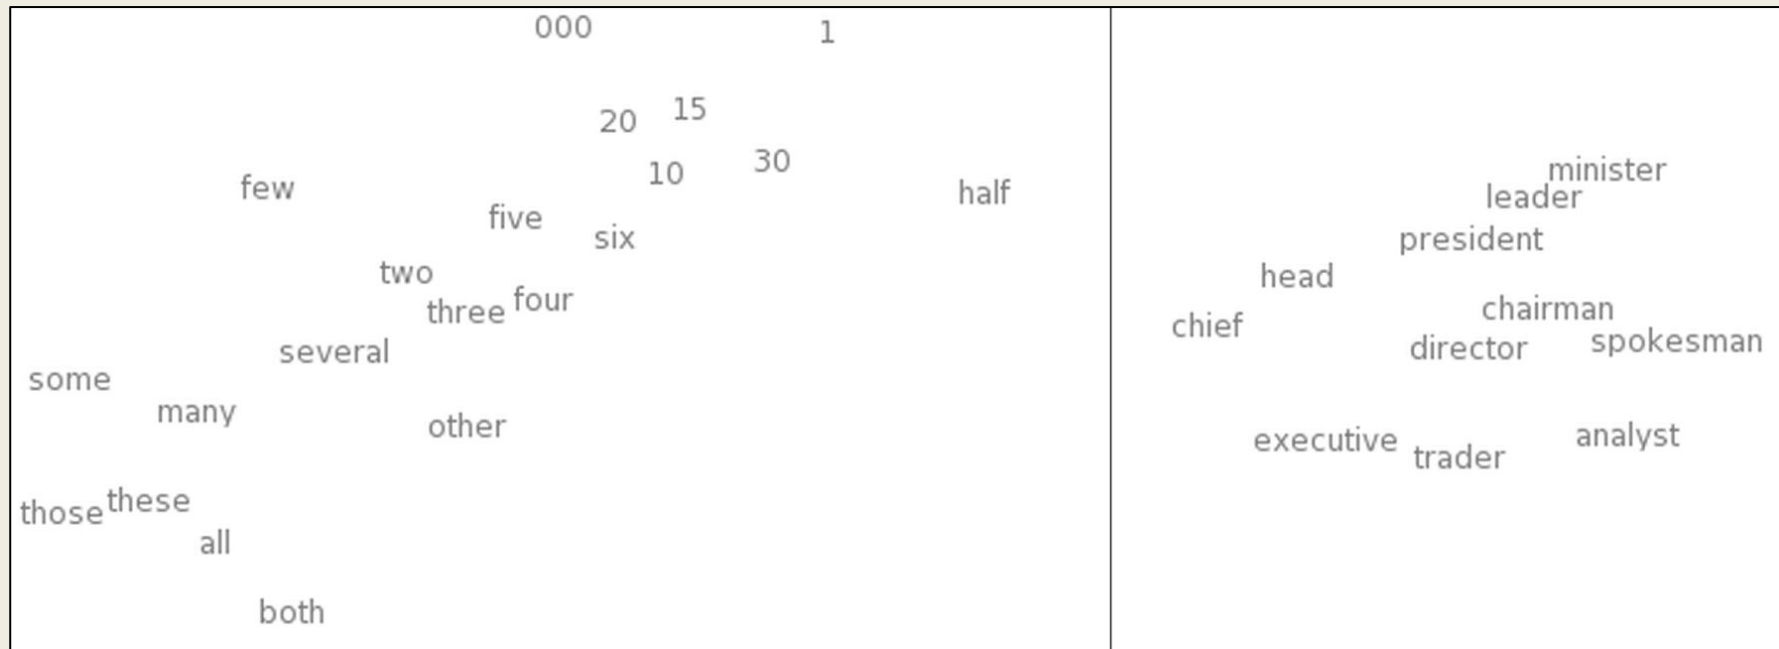
# Sentiment treebank





# Word embedding

Use vectors to represent words.



# NNs for parsing

NN + word embedding features:  
best parsers today.

Parser	UAS (%)	LAS (%)
MaltParser	90.93	88.95
Parsey McParseface	94.41	92.55

# Links

- <http://www.sciencefriction.net/blog/2010/12/06/1120/>
- <https://www.aclweb.org/anthology/E/E17/E17-2017.pdf>
- <https://research.googleblog.com/2016/05/announcing-syntaxnet-worlds-most.html>
- <http://www.cs.huji.ac.il/~danielh/P17-1104.pdf>
- <https://github.com/ayushoriginal/Sentiment-Analysis-Twitter>